

Chapter 2 - Files

2.1 What does a file need to run?

This chapter will be discussing how a file is loaded into memory and run as well as what is required for a file to run. We will be using the elf file format as an example.

2.2 ELF File Structure

2.2.1 Dissection of the ELF file structure

The ELF file, otherwise known as Executable and Linkable Format, is a common standard file format for executabl files, object code, shared libraries, and core dumps. This will be what you see most of the time when using linux.

In an elf file you will have the main ELF header, the program header table which describes zero or more memory segments, the section header table which describes zero or more sections, and then the data referred to by the tables.

ELF Header

Offset		Size in Bytes		Field	Purpose
32-bit	64-bit	32-bit	64-bit		
0x00		4		e_ident[EI_MAG0] through e_ident[EI_MAG3]	0x7F followed by ELF(45 4c 46) in ASCII; these four bytes constitute the magic number.
0x04		1		e_ident[EI_CLASS]	This byte is set to either 1 or 2 to signify 32- or 64-bit format, respectively.
0x05		1		e_ident[EI_DATA]	This byte is set to either 1 or 2 to signify little or big endianness, respectively. This affects interpretation of multi-byte fields starting with offset 0x10.
0x06		1		e_ident[EI_VERSION]	Set to 1 for the original and current version of ELF.
0x07		1		e_ident[EI_OSABI]	Identifies the target operating system application binary interface. The ABI is an interface between two binary program modules and defines how data structures or computational routines are accessed in machine code. Often set to 0 regardless of platform.

Offset		Size in Bytes		Field	Purpose																				
32-bit	64-bit	32-bit	64-bit																						
0x08		1		e_ident[EI_ABIVERSION]	Further specifies the ABI version. Its interpretation depends on the target ABI.																				
0x09		7		e_ident[EI_PAD]	Unused and filled with zeroes.																				
0x10		2		e_type	Identifies the object file type. <table><thead><tr><th>Value</th><th>Type</th></tr></thead><tbody><tr><td>0x00</td><td>ET_NONE</td></tr><tr><td>0x01</td><td>ET_REL</td></tr><tr><td>0x02</td><td>ET_EXEC</td></tr><tr><td>0x03</td><td>ET_DYN</td></tr><tr><td>0x04</td><td>ET_CORE</td></tr><tr><td>0xFE00</td><td>ET_LOOS</td></tr><tr><td>0xFEFF</td><td>ET_HIOS</td></tr><tr><td>0xFF00</td><td>ET_LOPROC</td></tr><tr><td>0xFFFF</td><td>ET_HIPROC</td></tr></tbody></table>	Value	Type	0x00	ET_NONE	0x01	ET_REL	0x02	ET_EXEC	0x03	ET_DYN	0x04	ET_CORE	0xFE00	ET_LOOS	0xFEFF	ET_HIOS	0xFF00	ET_LOPROC	0xFFFF	ET_HIPROC
Value	Type																								
0x00	ET_NONE																								
0x01	ET_REL																								
0x02	ET_EXEC																								
0x03	ET_DYN																								
0x04	ET_CORE																								
0xFE00	ET_LOOS																								
0xFEFF	ET_HIOS																								
0xFF00	ET_LOPROC																								
0xFFFF	ET_HIPROC																								
0x12		2		e_machine	Specifies target instruction set architecture.																				
0x14		4		e_version	Set to 1 for the original version of ELF.																				
0x18		4	8	e_entry	This is the memory address of the entry point from where the process starts executing.																				
0x1C	0x20	4	8	e_phoff	Points to the start of the program header table.																				
0x20	0x28	4	8	e_shoff	Points to the start of the section header table.																				
0x24	0x30	4		e_flags	Interpretation of this field depends on the target architecture.																				
0x28	0x34	2		e_ehsize	Contains the size of the header.																				
0x2A	0x36	2		e_phentsize	Contains the size of a program header table entry.																				
0x2C	0x38	2		e_phnum	Contains the number of entries in the program header table.																				
0x2E	0x3A	2		e_shentsize	Contains the size of a section header table entry.																				

Offset		Size in Bytes		Field	Purpose
32-bit	64-bit	32-bit	64-bit		
0x30	0x3C	2		e_shnum	Contains the number of entries in the section header table.
0x32	0x3E	2		e_shstrndx	Contains index of the section header table entry that contains the section names.
0x34	0x40				End of ELF Header.

Program header table

Offset		Size in Bytes		Field	Purpose	
32-bit	64-bit	32-bit	64-bit			
0x00		4	p_type	Identifies the type of the segment.		
				Value	Name	Meaning
				0x00000000	PT_NULL	Program header table entry unused.
				0x00000001	PT_LOAD	Loadable segment.
				0x00000002	PT_DYNAMIC	Dynamic linking information.
				0x00000003	PT_INTERP	Interpreter information.
				0x00000004	PT_NOTE	Auxiliary information.
				0x00000005	PT_SHLIB	Reserved.
				0x00000006	PT_PHDR	Segment containing program header table itself.
				0x00000007	PT_TLS	Thread-Local Storage template.
				0x60000000	PT_LOOS	Reserved inclusive range. Operating system specific.
				0x6FFFFFFF	PT_HIOS	
				0x70000000	PT_LOPROC	Reserved inclusive range. Processor specific.
				0x7FFFFFFF	PT_HIPROC	

Offset		Size in Bytes		Field	Purpose
32-bit	64-bit	32-bit	64-bit		
	0x04		4	p_flags	Segment-dependent flags.
0x04	0x08	4	8	p_offset	Offset of the segment in the file image.
0x08	0x10	4	8	p_vaddr	Virtual address of the segment in memory.
0x0C	0x18	4	8	p_paddr	On systems where physical address is relevant, reserved for segment's physical address.
0x10	0x20	4	8	p_filesz	Size in bytes of the segment in the file image. May be 0.
0x14	0x28	4	8	p_memsz	Size in bytes of the segment in memory. May be 0.
0x18		4		p_flags	Segment-dependent flags (position for 32-bit structure).
0x1C	0x30	4	8	p_align	0 and 1 specify no alignment. Otherwise should be a positive, integral power of 2, with p_vaddr equating p_offset modulus p_align.
0x20	0x38				End of Program Header.

Section Header table

Offset		Size in Bytes		Field	Purpose																					
32-bit	64-bit	32-bit	64-bit																							
0x00		4		sh_name	An offset to a string in the .shstrtab section that represents the name of this section.																					
0x04		4		sh_type	<div>Identifies the type of this header.</div> <table><thead><tr><th>Value</th><th>Name</th><th>Meaning</th></tr></thead><tbody><tr><td>0x0</td><td>SHT_NULL</td><td>Section header table entry unused</td></tr><tr><td>0x1</td><td>SHT_PROGBITS</td><td>Program data</td></tr><tr><td>0x2</td><td>SHT_SYMTAB</td><td>Symbol table</td></tr><tr><td>0x3</td><td>SHT_STRTAB</td><td>String table</td></tr><tr><td>0x4</td><td>SHT_RELA</td><td>Relocation entries with addends</td></tr><tr><td></td><td></td><td></td></tr></tbody></table>	Value	Name	Meaning	0x0	SHT_NULL	Section header table entry unused	0x1	SHT_PROGBITS	Program data	0x2	SHT_SYMTAB	Symbol table	0x3	SHT_STRTAB	String table	0x4	SHT_RELA	Relocation entries with addends			
Value	Name	Meaning																								
0x0	SHT_NULL	Section header table entry unused																								
0x1	SHT_PROGBITS	Program data																								
0x2	SHT_SYMTAB	Symbol table																								
0x3	SHT_STRTAB	String table																								
0x4	SHT_RELA	Relocation entries with addends																								

Offset		Size in Bytes		Field	Purpose		
32-bit	64-bit	32-bit	64-bit				
					0x5	SHT_HASH	Symbol hash table
					0x6	SHT_DYNAMIC	Dynamic linking information
					0x7	SHT_NOTE	Notes
					0x8	SHT_NOBITS	Program space with no data (bss)
					0x9	SHT_REL	Relocation entries, no addends
					0x0A	SHT_SHLIB	Reserved
					0x0B	SHT_DYNSYM	Dynamic linker symbol table
					0x0E	SHT_INIT_ARRAY	Array of constructors
					0x0F	SHT_FINI_ARRAY	Array of destructors
					0x10	SHT_PREINIT_ARRAY	Array of pre-constructors
					0x11	SHT_GROUP	Section group
					0x12	SHT_SYMTAB_SHNDX	Extended section indices
					0x13	SHT_NUM	Number of defined types.
					0x60000000	SHT_LOOS	Start OS-specific.
0x08		4	8	sh_flags	Identifies the attributes of the section.		
					Value	Name	Meaning
					0x1	SHF_WRITE	Writable
					0x2	SHF_ALLOC	Occupies memory during execution
					0x4	SHF_EXECINSTR	Executable
					0x10	SHF_MERGE	Might be

Offset		Size in Bytes		Field	Purpose	
32-bit	64-bit	32-bit	64-bit			
						merged
					0x20	SHF_STRINGS Contains null-terminated strings
					0x40	SHF_INFO_LINK 'sh_info' contains SHT index
					0x80	SHF_LINK_ORDER Preserve order after combining
					0x100	SHF_OS_NONCONFORMING Non-standard OS specific handling required
					0x200	SHF_GROUP Section is member of a group
					0x400	SHF_TLS Section hold thread-local data
					0x0ff00000	SHF_MASKOS OS-specific
					0xf0000000	SHF_MASKPROC Processor-specific
					0x4000000	SHF_ORDERED Special ordering requirement (Solaris)
					0x8000000	SHF_EXCLUDE Section is excluded unless referenced or allocated (Solaris)

Offset		Size in Bytes		Field	Purpose
32-bit	64-bit	32-bit	64-bit		
0x0C	0x10	4	8	sh_addr	Virtual address of the section in memory, for sections that are loaded.
0x10	0x18	4	8	sh_offset	Offset of the section in the file image.
0x14	0x20	4	8	sh_size	Size in bytes of the section in the file image. May be 0.
0x18	0x28	4	sh_link		Contains the section index of an associated section. This field is used for several purposes, depending on the type of section.
0x1C	0x2C	4	sh_info		Contains extra information about the section. This field is used for several purposes, depending on the type of section.
0x20	0x30	4	8	sh_addralign	Contains the required alignment of the section. This field must be a power of two.
0x24	0x38	4	8	sh_entsize	Contains the size, in bytes, of each entry, for sections that contain fixed-size entries. Otherwise, this field contains zero.
0x28	0x40				End of Section Header.

Section Headers

- .text
 - Contains executable code. It will be packed into a segment with read and execute access rights. It is only loaded once, as the contents will not change.
- .data
 - Initialized data, with read/write access rights.
- .rodata
 - Initialized data, with read access rights only.
- .bss
 - Uninitialized data, with read/write access rights.

Sections

- Symbol Table:

The symbol table is a section or sections that define the location, type, visibility, and other traits of various symbols declared in the source, created during compilation or linking, or otherwise present.

Each symbol table entry contains important information such as the symbol name, whether it is an external or internal symbol, the address of the segment (absolute or relative), and the symbol type and binding.

```
typedef struct elf32_sym{
    Elf32_Word    st_name;
    Elf32_Addr    st_value;
```

```

    Elf32_Word    st_size;
    unsigned char st_info;
    unsigned char st_other;
    Elf32_Half    st_shndx;
} Elf32_Sym;

typedef struct elf64_sym {
    Elf64_Word st_name;      /* Symbol name, index in string tbl */
    unsigned char st_info;   /* Type and binding attributes */
    unsigned char st_other;  /* No defined meaning, 0 */
    Elf64_Half st_shndx;     /* Associated section index */
    Elf64_Addr st_value;     /* Value of the symbol */
    Elf64_Xword st_size;     /* Associated symbol size */
} Elf64_Sym;

```

Bindings

STB_LOCAL	0
STB_GLOBAL	1
STB_WEAK	2
STB_LOOS	10
STB_HIOS	12
STB_LOPROC	13
STB_HIPROC	15

STB_LOCAL

Local symbol. These symbols are not visible outside the object file containing their definition. Local symbols of the same name can exist in multiple files without interfering with each other.

STB_GLOBAL

Global symbols. These symbols are visible to all object files being combined. One file's definition of a global symbol satisfies another file's undefined reference to the same global symbol.

STB_WEAK

Weak symbols. These symbols resemble global symbols, but their definitions have lower precedence.

STB_LOOS - STB_HIOS

Values in this inclusive range are reserved for operating system-specific semantics.

STB_LOPROC - STB_HIPROC

Values in this inclusive range are reserved for processor-specific semantics.

Types

STT_NOTYPE	0
STT_OBJECT	1
STT_FUNC	2
STT_SECTION	3
STT_FILE	4
STT_COMMON	5
STT_TLS	6
STT_LOOS	10
STT_HIOS	12
STT_LOPROC	13
STT_SPARC_REGISTER	13
STT_HIPROC	15

STT_NOTYPE

The symbol type is not specified.

STT_OBJECT

This symbol is associated with a data object, such as a variable, an array, and so forth.

STT_FUNC

This symbol is associated with a function or other executable code.

STT_SECTION

This symbol is associated with a section. Symbol table entries of this type exist primarily for relocation and normally have STB_LOCAL binding.

STT_FILE

Conventionally, the symbol's name gives the name of the source file that is associated with the object file. A file symbol has STB_LOCAL binding and a section index of SHN_ABS. This symbol, if present, precedes the other STB_LOCAL symbols for the file.

Symbol index 1 of the SHT_SYMTAB is an STT_FILE symbol representing the object file. Conventionally, this symbol is followed by the files STT_SECTION symbols. These section symbols are then followed by any global symbols that have been reduced to locals.

STT_COMMON

This symbol labels an uninitialized common block. This symbol is treated exactly the same as STT_OBJECT.

STT_TLS

The symbol specifies a thread-local storage entity. When defined, this symbol gives the assigned offset for the symbol, not the actual address.

Thread-local storage relocations can only reference symbols with type STT_TLS. A reference to a symbol of type STT_TLS from an allocatable section, can only be achieved by using special thread-local storage relocations. See Chapter 14, Thread-Local Storage for details. A reference to a symbol of type STT_TLS from a non-allocatable section does not have this restriction.

STT_LOOS - STT_HIOS

Values in this inclusive range are reserved for operating system-specific semantics.

STT_LOPROC - STT_HIPROC

Values in this inclusive range are reserved for processor-specific semantics.

- String Table

Consecutive zero terminated strings. There are normally several strings tables such as .strtab (default string table), .shstrtab (section string table), and .dynstr (string table for dynamic linking).

- BSS and SHT_NOBITS

The BSS is a block of memory which has been zeroed. This is an area in memory where variables with global lifetime that haven't been initialized are stored.

- Relocation Sections

Relocation starts with a table of relocation entries which can be used to locate relevant section headers. There are two kinds of relocation structures: one with explicit (SHT_RELA) and one without (SHT_REL). Info is used to compute the symbol the relocation applies to and what type of relocation should be applied.

```
typedef struct elf64_rel {
    Elf64_Addr r_offset; /* Location at which to apply the action */
    Elf64_Xword r_info; /* index and type of relocation */
} Elf64_Rel;

typedef struct elf64_rela {
    Elf64_Addr r_offset; /* Location at which to apply the action */
    Elf64_Xword r_info; /* index and type of relocation */
    Elf64_Sxword r_addend; /* Constant addend used to compute value */
} Elf64_Rela;
```

Relocation Types:

x64: Relocation Types

The relocations that are listed in the following table are defined for x64.

Table 12-16 x64: ELF Relocation Types

Name	Value	Field	Calculation
R_AMD64_NONE	0	None	None
R_AMD64_64	1	word64	$S + A$
R_AMD64_PC32	2	word32	$S + A - P$
R_AMD64_GOT32	3	word32	$G + A$
R_AMD64_PLT32	4	word32	$L + A - P$
R_AMD64_COPY	5	None	Refer to the explanation following this table.
R_AMD64_GLOB_DAT	6	word64	S
R_AMD64_JUMP_SLOT	7	word64	S
R_AMD64_RELATIVE	8	word64	$B + A$
R_AMD64_GOTPCREL	9	word32	$G + GOT + A - P$
R_AMD64_32	10	word32	$S + A$
R_AMD64_32S	11	word32	$S + A$
R_AMD64_16	12	word16	$S + A$
R_AMD64_PC16	13	word16	$S + A - P$
R_AMD64_8	14	word8	$S + A$
R_AMD64_PC8	15	word8	$S + A - P$
R_AMD64_PC64	24	word64	$S + A - P$
R_AMD64_GOTOFF64	25	word64	$S + A - GOT$
R_AMD64_GOTPC32	26	word32	$GOT + A + P$
R_AMD64_SIZE32	32	word32	$Z + A$
R_AMD64_SIZE64	33	word64	$Z + A$

Where:

A

The addend used to compute the value of the relocatable field.

B

The base address at which a shared object is loaded into memory during execution. Generally, a shared object file is built with a base virtual address of 0. However, the execution address of the shared object is different. See Program Header.

G

The offset into the global offset table at which the address of the relocation entry's symbol resides during execution. See Global Offset Table (Processor-Specific).

GOT

The address of the global offset table. See Global Offset Table (Processor-Specific).

L

The section offset or address of the procedure linkage table entry for a symbol. See Procedure Linkage Table (Processor-Specific).

P

The section offset or address of the storage unit being relocated, computed using `r_offset`.

S

The value of the symbol whose index resides in the relocation entry.

Z

The size of the symbol whose index resides in the relocation entry.